

Maya

SPEC · V3 · WHAT MAYA IS AND DOES

WHAT MAYA IS

Maya is the AI consultation engine for Mana Siyo. She listens as a client describes the garment they want, extracts the silhouette, aesthetic, and detail language from their words, lays it out as a moodboard, and renders three-view atelier drawings of the design — preserving the client's identity across every render so the model on the page is the actual person who'll wear the clothes.

WHO IT'S FOR

Fromsa (founder, Mana Siyo) uses Maya in real consultations with clients. The room is set up so the client speaks freely and Maya quietly captures, surfaces, and visualizes alongside.

WHAT IT DOES

CAPABILITY	WHAT IT DOES
Voice consultation	Live transcription (Web Speech) with Whisper fallback. Post-processed by gpt-4.1 to extract silhouette / aesthetic / detail chips.
Moodboard	Drag-drop inspo, paste files, dissect uploaded inspos into vocabulary chips. Pinterest-layout reorg on demand.
Visualize	Three-view atelier render (side · front · back) via gpt-image-2. Face photo + descriptors + lineage + style refs all feed into the multi-anchor /edits payload.
Modify / Re-cut	In-place edits — apply mods or swap fabric while preserving the model, pose, and construction pixel-for-pixel.
Favorites + Submit	Heart visions to a second screen; tap Submit on any favorite to save the full consultation folder (one-pager PDF, transcript, summary, moodboard.json, uploaded inspos).
Avatar drawer	Face photo + cached identity descriptors + body measurements. Drives every render's identity lock.
Fabrics drawer	In-house swatch library via File System Access API. Swatches feed Visualize and Re-cut.
One-pager export	Per-consultation PDF: hero image + James paragraph + vision text + measurements + selected inspos.
Parallel renders	Visualize / Modify / Re-cut all run concurrently. Status bar shows visualizing ×N.

ARCHITECTURE — ONE PARAGRAPH

Single HTML file (~7,800 lines). No build step. CSS in <style>, JS in <script>, both inlined. State persists via localStorage (API key, settings) and IndexedDB (items, face descriptors). OpenAI: gpt-image-2 for images, gpt-4.1 for chat and vision, gpt-4o-mini for fast helpers, whisper-1 for audio. File System Access API for the Clients / Inhouse Fabrics / Headshots folders.

THE TWO SCREENS

SCREEN	ROLE
Inspo (moodboard)	Voice bar bottom, MAYA wordmark + hamburger top, cards floating on the cosmic backdrop. The consultation surface.

SCREEN	ROLE
Favorites	Hearted visions only. Glowing glass cards in a horizontal scroller. Tap a card to open the submit-mode viewer; tap Submit to save the client folder.

THE DESIGN CANON — ONE LINE

Page 1 is the source of truth. Every other surface — splash, drawer, modals, viewer, toast, thank-you — uses the same liquid-glass pill recipe, the same Cormorant + Jost pairing, the same bone-white-over-cosmos palette. Pixel-level rules live in Aesthetics.pdf.

CURRENT VERSION

v10.2 — native scroll-snap on both axes (vertical screens + horizontal drawer). Submit saves to Clients folder. Canon-compliant chrome across every screen.